

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/45877649>

Engineering a Scalable High Quality Graph Partitioner

Article · October 2009

Source: arXiv

CITATIONS

20

READS

100

3 authors, including:



Manuel Holtgrewe

Freie Universität Berlin

13 PUBLICATIONS 280 CITATIONS

[SEE PROFILE](#)



Christian Schulz

University of Vienna

75 PUBLICATIONS 867 CITATIONS

[SEE PROFILE](#)

Some of the authors of this publication are also working on these related projects:



Kernelization [View project](#)



KaHyPar - Karlsruhe Hypergraph Partitioning [View project](#)

Engineering a Scalable High Quality Graph Partitioner

Manuel Holtgrew, Peter Sanders, Christian Schulz

Abstract

We describe an approach to parallel graph partitioning that scales to hundreds of processors and produces a high solution quality. For example, for many instances from Walshaw’s benchmark collection we improve the best known partitioning. We use the well known framework of multi-level graph partitioning. All components are implemented by scalable parallel algorithms. Quality improvements compared to previous systems are due to better prioritization of edges to be contracted, better approximation algorithms for identifying matchings, better local search heuristics, and perhaps most notably, a parallelization of the FM local search algorithm that works more locally than previous approaches.

1 Introduction

Many important applications of computer science involve processing large graphs, e.g., stemming from finite element methods, digital circuit design, route planning, social networks, . . . Very often these graphs need to be partitioned or clustered such that there are few edges between the blocks (pieces). In particular, when you process a graph in parallel on k PEs (processing elements) you often want to partition the graph into k blocks of about equal size. In this paper we focus on a version of the problem that constrains the maximum block size to $(1 + \epsilon)$ times the average block size and tries to minimize the total cut size, i.e., the number of edges that run between blocks. It is well known that there are more realistic (and more complicated) objective functions involving also the block that is worst and the number of its neighboring nodes [14] but minimizing the cut size has been adopted as a kind of standard since it is usually highly correlated with the other formulations. We believe that the results presented here will be adaptable to other objective functions and also to other setting such as graph clustering where k and the block sizes are not necessarily fixed.

We begin in Section 2 by introducing basic concepts. The main part of the paper are the sections on contraction 3, initial partitioning 4, and refinement 5. Section 6 summarizes extensive experiments done to tune the algorithm and evaluate its performance. Some related work is discussed in Section 7 and Section 8 summarizes the results and gives some outlook on future work.

2 Preliminaries

Consider an undirected graph $G = (V, E, c, \omega)$ with edge weights $\omega : E \rightarrow \mathbb{R}_{>0}$, node weights $c : V \rightarrow \mathbb{R}_{\geq 0}$, $n = |V|$, and $m = |E|$. We extend c and ω to sets, i.e., $c(V') := \sum_{v \in V'} c(v)$ and

$\omega(E') := \sum_{e \in E'} \omega(e)$. $\Gamma(v) := \{u : \{v, u\} \in E\}$ denotes the neighbors of v .

We are looking for *blocks* of nodes V_1, \dots, V_k that partition V , i.e., $V_1 \cup \dots \cup V_k = V$ and $V_i \cap V_j = \emptyset$ for $i \neq j$. The *balancing constraint* demands that $\forall i \in 1..k : c(V_i) \leq L_{\max} := (1 + \epsilon)c(V)/k + \max_{v \in V} c(v)$ for some parameter ϵ . The objective is to minimize the total *cut* $\sum_{i < j} w(E_{ij})$ where $E_{ij} := \{\{u, v\} \in E : u \in V_i, v \in V_j\}$. By default, our initial inputs will have unit edge and node weights. However, even those will be translated into weighted problems in the course of the algorithm.

A matching $M \subseteq E$ is a set of edges that do not share any common nodes, i.e., the graph (V, M) has maximum degree one.

An edge coloring \mathcal{C} assigns a color (a number) to each edge of a graph such that no two incident edges have the same color. Note that the edges with a particular color define a matching, i.e., \mathcal{C} partitions the edges into matchings. We will be interested in colorings with a small number of different colors used.

Contracting an edge $\{u, v\}$ means to replace the nodes u and v by a new node x connected to the former neighbors of u and v . We set $c(x) = c(u) + c(v)$. If replacing edges of the form $\{u, w\}, \{v, w\}$ would generate two parallel edges $\{x, w\}$, we insert a single edge with $w(\{x, w\}) = w(\{u, w\}) + w(\{v, w\})$. *Uncontracting* an edge e undos its contraction. In order to avoid tedious notation, G will denote the current state of the graph before and after a (un)contraction unless we explicitly want to refer to different states of the graph.

The multilevel approach to clustering consists of three main phases.

In the *contraction* (coarsening) phase, we iteratively identify matchings $M \subseteq E$ and contract the edges in M . This is repeated until $|V|$ falls below some threshold. Contraction should quickly reduce the size of the input and each computed level should reflect the global structure of the input network. In particular, nodes should represent densely connected subgraphs.

Contraction is stopped when the graph is small enough to be directly partitioned in the *initial partitioning phase* using some other algorithm. We could actually use a trivial initial partitioning algorithm if we contract until exactly k nodes are left. However, if $|V| \gg k$ we can afford to run some fairly expensive algorithm for initial partitioning.

In the *refinement* (or uncoarsening) phase, the matchings are iteratively uncontracted. After uncontracting a matching, the refinement algorithm moves nodes between blocks in order to reduce the cut size or balance. The nodes to move are often found using some kind of local search. The intuition behind this approach is that a good partition at one level of the hierarchy will also be a good partition on the next finer level so that refinement will quickly find a good solution.

3 Contraction

We distinguish two separate choices for computing a matching: A *rating function* for the edges telling us which edges are how valuable for the matching and a *matching algorithm* that tries to find a matching of near maximum weight efficiently. Contractions are run until the graph is “small enough”.

3.1 Edge Rating

In most previous work, the edge weight $\omega(e)$ itself is used as a rating function (see Section 7 for more details). We additionally consider

$$\begin{aligned}\text{expansion}(\{u, v\}) &:= \frac{\omega(\{u, v\})}{c(u) + c(v)} \\ \text{expansion}^*(\{u, v\}) &:= \frac{\omega(\{u, v\})}{c(u)c(v)} \\ \text{expansion}^{*2}(\{u, v\}) &:= \frac{\omega(\{u, v\})^2}{c(u)c(v)} \\ \text{innerOuter}(\{u, v\}) &:= \frac{\omega(\{u, v\})}{\text{Out}(v) + \text{Out}(u) - 2\omega(u, v)}\end{aligned}$$

where $\text{Out}(v) := \sum_{x \in \Gamma(v)} \omega(\{v, x\})$. These bounds are heuristically inferred from a few basic principles: it's good to contract heavy edges because this decreases the cut size. For the same reason we want to avoid clusters with many outgoing edges. Furthermore, we preferably contract light nodes because we want to keep the node weight at any level of contraction reasonably uniform.

In [15] several other functions based on ratings used in graph clustering are considered. However, they did not lead to very good results so that we do not go into details here.

3.2 Sequential Matching Algorithms

Although the maximum weight matching problem can be solved optimally in polynomial time, the available algorithms are too slow for very large graphs so that all graph partitioners use fast approximation algorithms. We tried three different matching algorithms that all run in linear or near linear time:

SHEM: *Sorted Heavy Edge Matching* is the algorithm used in Metis [22]. The nodes are sorted by increasing degree and then scanned. For each scanned node v , the heaviest edge $\{u, v\}$ incident to v is put into the matching and all remaining edges incident to u and v are excluded from further consideration. This algorithm is very fast but cannot give any worst case guarantees.

Greedy: The edges are sorted by descending weight and then scanned. When edge $\{u, v\}$ and neither u nor v are matched yet, $\{u, v\}$ is put into the matching. The Greedy algorithm guarantees a matching whose weight is at least half of the weight of a maximum weight matching.

GPA: The *Global Path Algorithm* was proposed in [17] as a synthesis of the Greedy algorithm and the Path Growing Algorithm [7]. All three algorithms achieve a half-approximation in the worst case, but empirically, GPA gives considerably better results. Similar to Greedy, GPA scans the edges in order of decreasing weight but rather than immediately building a matching, it first constructs a collection of paths and even cycles. Afterwards, optimal solutions are computed for each of these paths and cycles using dynamic programming.

We have not tried more sophisticated linear time algorithms that achieve 2/3-approximations since in [17] they empirically turn out to be much slower yet not much better than GPA.

3.3 Parallel Matching Algorithms

In our basic strategy we follow [16]. We first compute a preliminary partition of the graph, e.g., using coordinate information. Currently we have implemented a recursive bisection algorithm for nodes with 2D coordinates that alternately splits the data by the x -coordinate and the y -coordinate [2, 3]. We can also use the initial numbering of the nodes. Note that the initial partitioning does not directly affect the final partitioning computed later – its main purpose is to increase locality for the computation of matchings.

We then combine a sequential matching algorithm running on each partition and a parallel matching algorithm running on the *gap graph*. The gap graph consists of those edges $\{u, v\}$ where u and v reside on different PEs and $\omega(\{u, v\})$ exceeds the weight of the edges that may have been matched by the local matching algorithms to u and v . The parallel matching algorithm itself iteratively matches edges that $\{u, v\}$ are locally heaviest both at u and v until no more edges can be matched.

4 Initial Partitioning

The contraction is stopped when the number of remaining nodes on some PE is below $\max(20, n/(\alpha k^2))$ for some tuning parameter α . The graph is then small enough to be partitioned on a single PE. Our framework allows using pMetis or Scotch for initial partitioning. We use the sequential algorithms and run them simultaneously on all PEs, each with a different seed for the random number generator. Since initial partitioning is very fast, it is also repeated several times. The best solution is then broadcast to all PEs.

5 Refinement

Recall that the refinement phase iteratively uncontracts the matchings contracted during the contraction phase. After a matching is uncontracted, local search based refinement algorithms move

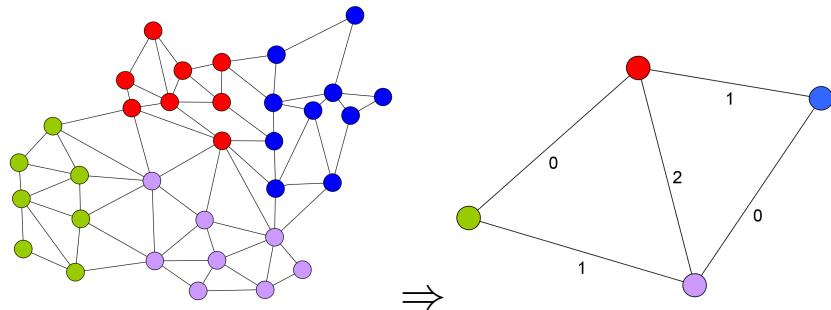


Figure 1: A graph which is partitioned into four blocks and its corresponding quotient graph \mathcal{Q} . The quotient graph has an edge coloring indicated by the numbers and each edge set induced by edges with the same color form a matching $\mathcal{M}(c)$. Pairs of blocks with the same color can be refined in parallel.

nodes between block boundaries in order to reduce the cut while maintaining the balancing constraint. As most other current systems, we adopt the basic approach from [10] which runs in linear time. The basic idea behind our parallel refinement algorithm is that at any time, each PE may work on one pair of neighboring blocks performing a local search constrained to moving nodes between these two blocks. In order to assign pairs of blocks to PEs, we use the *quotient graph* Q whose nodes are blocks of the current partition and whose edges indicate that there are edges between these blocks in the underlying graph G . Since we have the same number of PEs and blocks, each PE will work the block assigned to it and at one of its neighbors in Q . From now on, we will therefore identify blocks and PEs. Figure 1 gives an example.

We use matchings of Q to define with which neighbor in Q a PE is working at a particular point in time. If u, v is in the matching, both corresponding PEs will refine the partitions u and v using different seeds for their random number generator. See Section 5.2 for more details. After the local search is finished, the better partitioning of the two blocks is adopted.

Of course, for a good partition, we need to perform local search on every edge of Q eventually (we call this a *global iteration*). Section 5.1 describes our approaches for ensuring this.

We ensure this by iterating through the matchings defined by an edge coloring of Q . See Section 5.1 for more details.

Overall, this approach naturally defines a nested loop controlling our local search strategy. The innermost loop moves nodes between two blocks using the FM-algorithm [10]. A *local iteration* repeats this local search. A *global iteration* iterates over the colors of an edge coloring. The loops terminate when either no improvement was found (in strong variants: when no improvement was found twice in a row.) or when a preset maximum number of iterations is exceeded.

5.1 Choosing Matchings

We have implemented two strategies. One finds edges of Q not yet used for local search in a randomized local way. The other steps through the colors of an edge coloring of the quotient graph Q . Note that this requires only local synchronization between PEs actually collaborating at a particular point in time. We only describe the latter one here since it performs slightly better in our experiments. Our coloring algorithm is a parallelization of a well known sequential greedy edge coloring algorithm: Each PE has a set \mathcal{L} of free colors that have not been used for coloring incident edges. In each round of the algorithm, PEs throw a coin with sides *active* and *passive*. An active PE u picks a random incident uncolored edge $\{u, v\}$ and sends this edge together with its free-list to PE v . These *requests* are rejected if they are sent to other active PEs. Passive PEs v process requests $(\{u, v\}, \mathcal{L}')$ by choosing the color $c = \min L \cap L'$ for edge $\{u, v\}$ and sending c back to u . This algorithm is repeated until all edges are colored. It can be shown that this algorithm needs at most twice as many colors as an optimal edge coloring.

5.2 Refinement Between Two Blocks

We use a fully distributed graph data structure. More precisely, we use hybrid between a static and a dynamic graph data structure. Immediately after uncontracting a matching, every PE stores the partition it is responsible for in a static adjacency array representation (also called forward-star



Figure 2: Refinement between two blocks using boundary exchange.

representation), i.e., there is an edge array storing target nodes and edge weights and a node array storing node weights and the start of the relevant segment in the edge array. In addition, we use a hash table to store migrated nodes and a second edge array for the corresponding edges. See [23] for more details. Before a local search operation, we perform a bounded breadth first search starting from the boundary of each block, and send copies of this boundary array to the partner PE in the local search. The local search is then limited to this boundary area. This way, for large graphs, only a small fraction of each block has to be communicated. If it should really happen that the local search would profit from going beyond the boundary area, this will be possible in the next iteration of some of the outer loops. Figure 2 shows this schematically.

The local search algorithm itself is basically the FM-algorithm [10]: For each of the two blocks A, B under consideration, a PE keeps a priority queue of nodes eligible to move. The priority is based on the *gain*, i.e., the decrease in edge cut when the node is moved to the other side. Each node is moved at most once within a single local search. The queues are initialized in random order with the nodes at the partition boundary. We have tried several queue selection strategies: *Alternating* between A and B [10], *MaxLoad* where always the heavier block gives a node, and *TopGain*, where the queue promising larger gain is used. In order to achieve a good balance, *TopGain* adopts the exception that *MaxLoad* is used when one of the blocks is overloaded. When not otherwise mentioned, we use *TopGain* with random tie breaking. There is also a variant *TopGainMaxLoad* that uses *MaxLoad* when both queues promise the same gain.

The search is broken when more than $\alpha \min\{|A|, |B|\}$ nodes have been moved without yielding an improvement. When the search stops, search is rolled back to the state with the lexicographically best value of the tuple $(imbalance, cutValue)$. Where $imbalance$ is $\max(0, \max(c(A) - L_{\max}, c(B) - L_{\max}))$.

6 Experiments

Implementation. We have implemented the algorithm described above using C++ and MPI. Overall, our program consists of about 34 000 lines of code. Priority queues for the local search are based on binary heaps. Hash tables use the library (extended STL) provided with the GCC compiler.

System. We have run our code on cluster with 200 nodes each equipped with two Quad-core Intel Xeon processors (X5355) which run at a clock speed of 2.667 GHz, have 2x4 MB of level 2 cache each and run Suse Linux Enterprise 10 SP 1. All nodes are attached to an InfiniBand 4X DDR interconnect which is characterized by its very low latency of below 2 microseconds and a point to point bandwidth between two nodes of more than 1300 MB/s. Our program was compiled

using GCC Version 4.3.1 and optimization level 3 using OpenMPI 1.2.8. Henceforth, a PE is one core of this machine.

Instances. We report experiments on two suites of instances summarized in Table 1. *rggX* is a *random geometric graph* with 2^X nodes where nodes represent random points in the unit square and edges connect nodes whose Euclidean distance is below $0.55\sqrt{\ln n/n}$. This threshold was chosen in order to ensure that the graph is almost connected. *DelaunayX* is the Delaunay triangulation of 2^X random points in the unit square. Graphs *bcsstk29..fetooth* and *ferotor..auto* come from Chris Walshaw’s benchmark archive [28]. Graphs *bel*, *nld*, *deu* and *eur* are undirected versions of the road networks of Belgium, the Netherlands, Germany, and Western Europe respectively, used in [5]. Instances *af_shell9* and *af_shell10* come from the Florida Sparse Matrix Collection [4]. *coAuthorsDBLP*, *citationCiteseer* are examples of social networks taken from [12]. Coordinate information is available for *rggX*, *DelaunayX*, the road networks, *bel*, *nld*, *deu* and *eur*, and for the finite element graphs *fecocean* and *fetooth*.

For the number of partitions k we choose the values used in [28]: 2, 4, 8, 16, 32, 64. Our default value for the allowed imbalance is 3 % since this is one of the values used in [28] and the

graph	n	m	graph	n	m
rgg17	2^{17}	1 457 506	rgg20	2^{20}	13 783 240
rgg18	2^{18}	3 094 566	Delaunay20	2^{20}	12 582 744
Delaunay17	2^{17}	786 352	fetooth	78 136	905 182
Delaunay18	2^{18}	1 572 792	598a	110 971	1 483 868
bcsstk29	13 992	605 496	ocean	143 437	819 186
4elt	15 606	91 756	144	144 649	2 148 786
fesphere	16 386	98 304	wave	156 317	2 118 662
cti	16 840	96 464	m14b	214 765	3 358 036
memplus	17 758	108 384	auto	448 695	6 629 222
cs4	33 499	87 716	deu	4 378 446	10 967 174
pwt	36 519	289 588	eur	18 029 721	44 435 372
bcsstk32	44 609	1 970 092	af_shell10	1 508 065	51 164 260
body	45 087	327 468	coAuthorsDBLP	299 067	1 955 352
t60k	60 005	178 880	citationCiteseer	434 102	32 073 440
wing	62 032	243 088			
finan512	74 752	522 240			
ferotor	99 617	1 324 862			
bel	463 514	1 183 764			
nld	893 041	2 279 080			
af_shell9	504 855	17 084 020			

Table 1: Basic properties of the graphs from our benchmark set. left: small to medium sized inputs, right: large instances. The latter class is split into five groups: geometric graphs, FEM graphs, street networks, sparse matrices, and social networks. Within their groups, the graphs are sorted by size.

default value in Metis.

When not otherwise mentioned, we perform 10 repetitions of each run and report the average result. When averaging over multiple instances, we use the geometric mean in order to give every instance the same influence on the final figure.

6.1 Configuring the Algorithm

Any multilevel algorithm has a considerable number of choices between algorithmic components and tuning parameters. In the following we explore the most important of these choices. In each case we will infer either a single “good” setting or two choices: the *fast* setting aims at a low execution time that still gives good partitioning quality and the *strong* setting targets good partitioning quality without investing an outrageous amount of time. At no point we tune parameters specifically for one instance. All other parameters are fixed at the default choices. When not otherwise mentioned, we use the *fast* parameter setting. For some of the values we do not show experiments to save space and because the experiments we did try do not give much new insight. Table 2 summarizes the settings. There is also a *minimal* variant where for all parameters the smallest possible value is chosen. Although the minimal variant can be viewed as overly crippled, it is useful when comparing to other, faster solvers.

Edge Ratings. Table 3 shows the average performance for different edge ratings. Note that the plain edge weight is considerably worse than the other ratings – up to 8.8 %. The other ratings are fairly close to each other and further experiments indicate that the remaining differences heavily depend on the instances and other parameters of the strategy. We adopt expansion^{*2} in the following.

parameter	minimal	fast	strong
rating		expansion ^{*2}	
matching		GPA	
stop contraction		$n/60k^2$	
init. part.		Scotch	
init. repeats	1	3	5
queue selection		TopGain	
BFS search depth	1	5	20
stop refinement	-	no change	2× no change
max. global iterations	1	15	15
local iterations	1	3	5
matching selection		distr. edge coloring	
FM-patience α	1 %	5 %	20 %
avg. cut (geom.)	2985	2910	2890
avg. time (geom.)[s]	0.67	1.29	2.10

Table 2: Parameter settings the for our main strategies.

Edge Rating	avg.	best.	avg. bal.	avg. t	Seq. Match.	avg.	best.	avg. bal.	avg. t
expansion*2	2910	2819	1.025	1.29	gpa	2910	2819	1.025	1.29
expansion*	2914	2815	1.025	1.30	shem	2984	2883	1.025	1.29
innerOuter	2914	2816	1.025	1.32	greedy	3854	3347	1.025	1.78
expansion	2940	2841	1.025	1.31					
weight	3165	3010	1.026	1.40					

Table 3: Results for KaPPa-Fast for different edge ratings and matching algorithms.

Sequential Matching Algorithm. In Table 3, we see that the other algorithms have at least 2.5 % worse edge cuts than GPA. Note that the overall running time in both configurations is about the same – although GPA is slower than SHEM, this disadvantage is offset by less work in the refinement phase. The greedy algorithm performs worse than the other strategies. This is astonishing since in [17] it produces fairly good results. Moreover, in the sequential experiments in [15] it also works well and outperforms SHEM. Apparently, there are some negative interactions with the parallelization here.

Initial Partitioning. So far, we tried pMetis and Scotch for initial partitioning. pMetis is about 4.7 % worse than Scotch and only has slightly lower overall runtime. We therefore adopt it as our default initial partitioner.

Queue Selection. Table 4 indicates that TopGain gives about 3.2 % better solutions than the more standard MaxLoad strategy. Interestingly, the details of the strategy are very important. Without resorting to MaxLoad in an overloaded situation we would not be able to fulfill the balance constraint. On the other hand, even using MaxLoad for tie breaking we are already worse than the seemingly stupid Alternating rule.

Global Iterations, Local Iterations, BFS Depth, and Local Search Parameters. For these parameters we get the predictable effect that more work yields better solutions albeit at a decreasing return on investment. It is then hard to say what parameters would be optimal. Roughly, our fast strategy represents values that yield execution times no more than 20 % larger than for the smallest possible value. These increases in execution time add up to 63 % more execution time than the fast strategy on average.

Variant	avg.	best.	bal.	avg. t
KaPPa-Strong	24227	23739	1.028	36.93
KaPPa-Fast	24725	24254	1.028	21.40
KaPPa-Minimal	26720	26005	1.028	5.94
seq. scotch	26811	-	1.027	5.95
kmetis	28705	26904	1.026	0.79
parmetis	31523	30449	1.041	0.59

Table 4: Left: Results for KaPPa-Fast for different queue selection strategies. Right: Comparison with other tools.

6.2 Comparison with other Partitioners

We now switch to our suite of larger graphs since that's what KaPPa was designed for and because we thus avoid the effect of over-tuning our algorithm parameters to the instances used for calibration.

Table 4 compares the performances of KaPPa with Scotch, kMetis (sequential) and parMetis (parallel). Detailed, per instance results can be found in Appendix A. parMetis produces about 30 % larger cuts than the strong variant of KaPPa, 27 % more than the fast one, and still 18 % more than the minimal one. Note that these differences are much larger than what can be obtained by just repeated runs, which gives only about 3 % improvement for 10 repetitions. Moreover parMetis is not able to fully adhere to the balancing constraint. On the other hand, parMetis is at least an order of magnitude faster.

For kMetis the differences are 18 %, 16 % and 7% respectively. For Scotch, we get 10 % for the strong variant, 8 % for the fast variant, and similar partitioning quality as for the weak variant. Comparing average execution times of parallel KaPPa with the sequential algorithms scotch and kMetis makes little sense because this depends a lot on the number of PEs used.

Although a large gap between the running times remains, the differences get smaller if one only considers graphs for which the current implementation of KaPPa was optimized: large graphs with coordinate information that allows geometric prepartitioning. Table 5 in the appendix shows data for the four graphs in our benchmark suite that have at least one million nodes and coordinate information (*rgg20*, *Delaunay20*, *deu*, *eur*). First note that for the European road network, *eur*, KaPPa produces a several times smaller cut than Metis. Apparently, Metis was not able at all to discover the structure inherent in the network (e.g., due to waterbodies, mountains, and national borders). KaPPa-minimal now outperforms Scotch, comes close to kMetis and is only a factor 3–6 slower than parMetis. Also note that the absolute execution times are in the range of a few seconds – few applications working on such large graphs will work on that time scale. Another interesting observation is that none of the other algorithms consistently complies with the balance constraint of 3 %. This is astonishing since these graphs have a very “harmless” structure – they are near planar (except for *rgg*) and have low maximum degree). It seems that our approach of careful, pairwise refinement successfully avoids such problems.

For the largest graphs available to us, we have scaled the number of processors further up to 1024. In Figure 3 we see that KaPPa¹ scales well all the way to the largest number of processors, while parMetis reaches its limit of scalability at around 100 PEs. Eventually, parMetis is slower than the fastest variant of KaPPa.

6.3 The Walshaw Benchmark

We now apply KaPPa to Walshaw's benchmark archive [28, 24] using the rules used there, i.e., running time is no issue but we want to achieve minimal cut values for $k \in \{2, 4, 8, 16, 32, 64\}$ and balance parameter $\epsilon \in \{0.01, 0.03, 0.05\}$. Thus, we further strengthen the strong strategy: We try each of the edge ratings innerOuter, expansion*, and expansion*² 50 times; BFS search depth is 20;

¹The minimal variant scales up to 512 PEs but this could be repaired by breaking the contraction later.

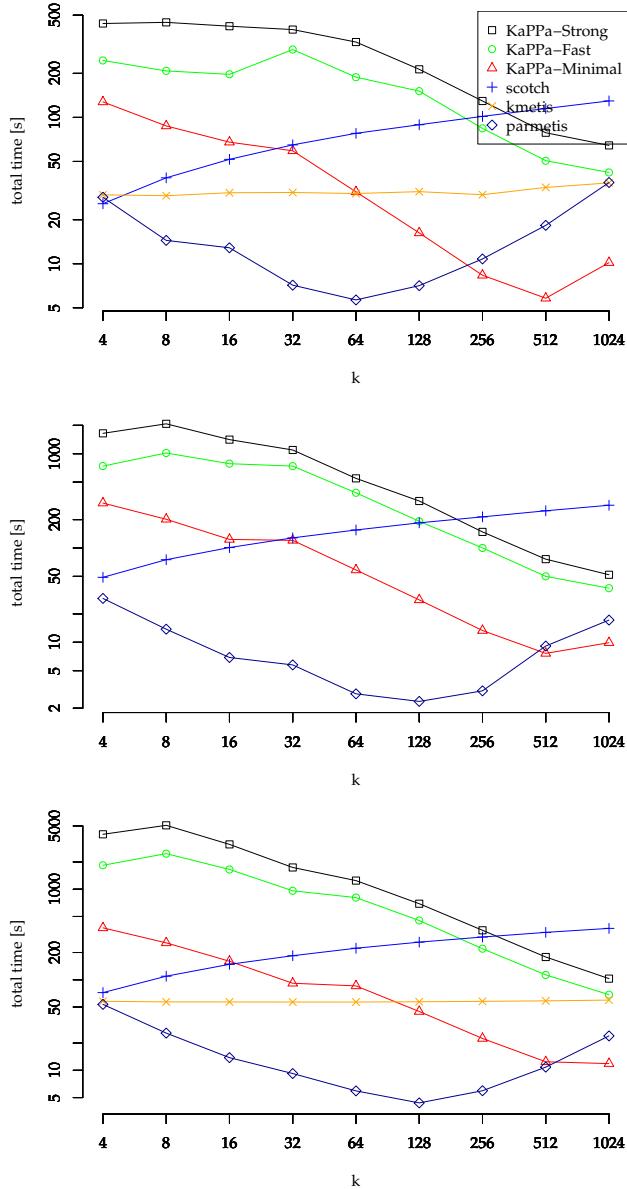


Figure 3: Scalability for graphs eur, rgg25, and Delaunay25.

FM patience $\alpha = 30\%$. Tables 21–23 in the Appendix show the results (left: KaPPa, right: best previous value) indicating an edge rating function that achieved our result. We obtain 54 improved entries for balance 5 %, 46 improvements for 3 %, and 31 improvements for balance 1 %. One interpretation is that the improvement due to the TopGain queue selection strategy become less effective for very small imbalance. Indeed, for balance 0 TopGain yields no improvements.² For

²However, the MaxLoad strategy given some slack on the balance constraint, yields good solutions that, for small k , are often fully balanced and yield improved values.

11 out of 14 instances from the large graphs we obtain improvements somewhere and for 9 out of 20 small instances (for all but two of the small instances we sometimes find a solution with the best known cut). The biggest absolute improvement is observed for instance *add32* at 1 % imbalance, and $k = 64$ where the old partition cuts 45 % more edges. We obtain few improvements for $k = 2$, perhaps still lacking specialized techniques for that case. We have many improvements for $k = 4$ going down for smaller graphs and larger k . Perhaps this could be changed by combining KaPPa with evolutionary techniques such as [24]. For large k we expect evolutionary methods to be superior to plain restarts that then have trouble exploring a sufficient part of the solution space.

7 Related Work

This paper is a summary and extension of the diploma theses [23, 15]. There has been a huge amount of research on graph partitioning so that we refer to overview papers such as [11, 22, 27] for a general overview. From now on focus on issues closely related to the contributions of our paper. All successful methods that are able to obtain good partitions for large real world graphs are based on the multilevel principle outlined in Section 2. The basic idea can be traced back to multigrid solvers for solving systems of linear equations [25, 9] but more recent practical methods are based on mostly graph theoretic aspects in particular edge contraction and local search. Well known software packages based on this approach include Chaco [13], Jostle [27], Metis [22], Party [8], and Scotch [19]. While Chaco and Party are no longer developed and have no parallel version, the others have been parallelized also. Probably the fastest available parallel code is the parallel version of Metis, parMetis. However, its partitioning quality is worse than the sequential version kMetis. In general it seems to be the case that previous parallelizations came with a penalty in partitioning quality. In contrast, our parallelization approach seems to *improve* partitioning quality.

The parallel version of Jostle [27] is similar to our approach since it applies local search to pairs of neighboring partitions. However, this parallelization has problems maintaining the balance of the partitions since at any particular time, it is difficult to say how many nodes are assigned to a particular block. We solve this problems by performing concurrent local searches only on independent pairs of partitions.

PT-Scotch, the parallel version of Scotch is based on recursive bipartitioning. This is more difficult to parallelize than direct k -partitioning since in the initial bipartition, there is less parallelism available. The unused processor power is used by performing several independent attempts in parallel. The involved communication effort is reduced by considering only nodes close to boundary of the current partitioning (band-refinement). We also use band-refinement but using a different algorithm and with much less replication of work.

DiBaP [18] is a multi-level graph partitioning package based on diffusion. It currently yields the best partitioning results for the biggest graphs in [26] but has no scalable parallelization.

Most previous approaches use the edge weight to quantify with which preference it is included into a matching. In [1], many different edge ratings are considered. However all of them use a very simple rating as the primary sorting criterion. In contrast, our approach genuinely combines the two sometimes conflicting criteria of contracting heavy edges and light vertices.

The need for fast, (near) linear time algorithms for approximate weighted matchings in

hierarchical graph partitioning has been a major motivation for developing such algorithms [21, 7, 6, 20, 17]. In contrast to the heavy edge matching algorithms used in most systems, these schemes give approximation guarantees of 1/2 [21, 7] or 2/3 [6, 20]. In [17] we developed another 1/2 algorithm that turned out to be even better than the 2/3 algorithms in many practical cases. Interestingly, only few of these results have so far found their way into actual graph partitioners. One contribution of our paper is to try them out.

8 Conclusions and Future Work

We have demonstrated that high quality graph partitioning can be done in parallel in a scalable way. This success is due to several innovations/observations that might also work in the framework of other graph partitioning and graph clustering systems: Edge rating functions that take into account other aspects than edge weight give considerably better results (8.8 % on the average for the experiments in Section 6.1). In particular, it seems that discouraging heavy nodes leads to much more uniform contraction all over the graph. High quality matching algorithms like GPA also yield a few percent improvement. In particular, the computational overhead for these algorithms is not affecting the overall runtime of a high quality graph partitioner, presumably because of less work in the refinement phase. FM-style local search can also yield improved quality if the highest gain queue is selected if possible. Feasibility can be maintained using an exception for overloaded blocks. Again, a few percent improvement in solution quality can be obtained. Perhaps the most surprising result is that localizing the local search to two blocks at a time does at the same time enable parallelization and *improve* partitioning quality compared to global local search. Although the individual improvement due to each improvement is relatively small, they add up to a sizable overall improvement. Also note that within a less tuned system, adding one of the improvements may have a larger effect than in a code with all improvements at once.

The current implementation of KaPPa is a research prototype rather than a widely usable tool. But considering its good results, we want to further improve it and advance it into a fully usable system usable for all kinds of inputs ranging from small graphs better handled by a lean sequential implementation to huge graphs with billions of nodes.

Besides many implementation issues that will hopefully improve execution time, the main conceptual task will be a generalization of the interface. We want a system where the number of partitioning PEs P and the number of blocks k can be chosen independently. This is rather straight forward when $k > P$ since this actually increases the amount of parallelism. For $k < P$, we could simply assign more PEs to the same local search (using different seeds). This would improve quality but reduces scalability and will not work for huge graphs where block sizes may exceed local memory size. Therefore, we need a parallel refinement algorithm working on only two neighbouring blocks. We also want to improve performance for graphs that are neither prepartitioned nor equipped with coordinates. The easiest solution for moderate P will be to use parMetis for initial partitioning. For very large systems we want to develop a very fast prepartitioner that works purely graph theoretically. A core component will be fast scalable parallel contraction. There will also be further issues when KaPPa is generalized for graph clustering, hypergraph partitioning, or repartitioning. Besides improving the functionality of KaPPa, there are also many ways to improve its

basic performance. In particular, it would be desirable to implement a more efficient representation of the distributed graph data structure.

Besides improving functionality of KaPPa, many interesting research questions remain. For example, one should investigate rating functions for edge contraction more systematically. Other refinement algorithms, e.g., based on flows or diffusion could be tried within our framework of pairwise refinement.

References

- [1] Amine Abou-Rjeili and George Karypis. Multilevel algorithms for partitioning power-law graphs. In *International Parallel & Distributed Processing Symposium*, 2006.
- [2] J. L. Bentley. Multidimensional binary search trees used for associative searching. *Commun. ACM*, 18(9):509–517, 1975.
- [3] M. J. Berger and S. H. Bokhari. A partitioning strategy for pdes across multiprocessors. In *ICPP*, pages 166–170, 1985.
- [4] T. Davis. The University of Florida Sparse Matrix Collection, <http://www.cise.ufl.edu/research/sparse/matrices>, 2008.
- [5] D. Delling, P. Sanders, D. Schultes, and D. Wagner. Engineering route planning algorithms. submitted for publication, <http://il1www.ira.uka.de/extra/publications/dssw-erpa-09.pdf>, 2008.
- [6] D. Drake and S. Hougardy. Improved linear time approximation algorithms for weighted matchings. In *7th International Workshop on Randomization and Approximation Techniques in Computer Science (APPROX)*, LNCS 2764, pages 14–23, 2003.
- [7] D. Drake and S. Hougardy. A simple approximation algorithm for the weighted matching problem. *Information Processing Letters*, 85:211–213, 2003.
- [8] R. Preis et al. PARTY partitioning library. <http://wwwcs.uni-paderborn.de/fachbereich/AG/monien/RESEARCH/PART/party.html>.
- [9] R. P. Fedorenko. A relaxation method for solving elliptic difference equations. *USSR Comput. Math. and Math. Phys.*, 5(1):1092–1096, 1961.
- [10] C. M. Fiduccia and R. M. Mattheyses. A Linear-Time Heuristic for Improving Network Partitions. In *19th Conference on Design Automation*, pages 175–181, 1982.
- [11] P.O. Fjallstrom. Algorithms for graph partitioning: A survey. *Linkoping Electronic Articles in Computer and Information Science*, 3(10), 1998.
- [12] R. Geisberger, P. Sanders, and D. Schultes. Better approximation of betweenness centrality. In *10th Workshop on Algorithm Engineering and Experimentation*, pages 90–108, San Francisco, 2008. SIAM.
- [13] B. Hendrickson. Chaco: Software for partitioning graphs. <http://www.sandia.gov/~bahendr/chaco.html>.

- [14] B. Hendrickson. Graph partitioning and parallel solvers: Has the emperor no clother? (extended abstract). In *IRREGULAR*, pages 218–225, 1998.
- [15] M. Holtgrewe. A scalable coarsening phase for a multi-level partitioning algorithm. Diploma thesis, Universität Karlsruhe, 2009.
- [16] F. Manne and R. H. Bisseling. A parallel approximation algorithm for the weighted maximum matching problem. In *7th International Conference on Parallel Processing and Applied Mathematics (PPAM)*, volume 4967 of *LNCS*, pages 708–717. Springer, 2007.
- [17] J. Maue and P. Sanders. Engineering algorithms for approximate weighted matching. In *6th International Workshop on Experimental Algorithms (WEA)*, volume 4525 of *LNCS*, pages 242–255. Springer, 2007.
- [18] H. Meyerhenke, B. Monien, and T. Sauerwald. A new diffusion-based multilevel algorithm for computing graph partitions of very high quality. In *IEEE International Symposium on Parallel and Distributed Processing, 2008. IPDPS 2008.*, pages 1–13, 2008.
- [19] F. Pellegrini. Scotch home page. <http://www.labri.fr/pelegrin/scotch>.
- [20] S. Pettie and P. Sanders. A simpler linear time $2/3 - \epsilon$ approximation for maximum weight matching. *Information Processing Letters*, 91(6):271–276, 2004.
- [21] R. Preis. Linear time $1/2$ -approximation algorithm for maximum weighted matching in general graphs. In *Proc. 16th Ann. Symp. on Theoretical Aspects of Computer Science (STACS)*, *LNCS 1563*, pages 259–269, 1999.
- [22] K. Schloegel, G. Karypis, and V. Kumar. Graph partitioning for high performance scientific simulations. In J. Dongarra et al., editor, *CRPC Parallel Computing Handbook*. Morgan Kaufmann, 2000. <http://www-users.cs.umn.edu/~karypis/publications/partitioning.html>.
- [23] C. Schulz. Scalable parallel refinement of graph partitions. Diploma thesis, Universität Karlsruhe, 2009.
- [24] A. J. Soper, C. Walshaw, and M. Cross. A combined evolutionary search and multilevel optimisation approach to graph partitioning. *J. Global Optimization*, 29(2):225–241, 2004.
- [25] R. V. Southwell. Stress-calculation in frameworks by the method of “Systematic relaxation of constraints”. *Proc. Roy. Soc. Edinburgh Sect. A*, pages 57–91, 1935.
- [26] C. Walshaw. The Graph Partitioning Archive, <http://staffweb.cms.gre.ac.uk/~c.walshaw/partition/>, 2008.
- [27] C. Walshaw and M. Cross. JOSTLE: Parallel Multilevel Graph-Partitioning Software – An Overview. In F. Magoules, editor, *Mesh Partitioning Techniques and Domain Decomposition Techniques*, pages 27–58. Civil-Comp Ltd., 2007. (Invited chapter).
- [28] C. Walshaw, M. Cross, Centre for Numerical Modelling, Process Analysis, and University of Greenwich. Mesh Partitioning: A Multilevel Balancing and Refinement Algorithm. *SIAM Journal on Scientific Computing*, 22(1):63–80, 2000.

A Detailed Results for the Large Instances.

alg.	k	graph	avg. cut	best. cut.	avg. balance	avg. runtime
KaPPa-strong	64	rgg20	35354	34778	1.030	11.62
KaPPa-strong	64	Delaunay20	25179	24799	1.030	22.04
KaPPa-strong	64	deu	4093	4021	1.029	49.55
KaPPa-strong	64	eur	5393	5290	1.030	308.17
KaPPa-fast	64	rgg20	35539	35086	1.030	9.95
KaPPa-fast	64	Delaunay20	25129	24946	1.030	12.83
KaPPa-fast	64	deu	4146	4078	1.029	31.63
KaPPa-fast	64	eur	5538	5448	1.030	183.98
KaPPa-minimal	64	rgg20	35629	35252	1.030	2.09
KaPPa-minimal	64	Delaunay20	27001	26314	1.029	1.79
KaPPa-minimal	64	deu	4317	4193	1.029	5.97
KaPPa-minimal	64	eur	5770	5569	1.029	29.64
Scotch	64	rgg20	38815	38815	1.031	9.84
Scotch	64	Delaunay20	26163	26163	1.037	7.36
Scotch	64	deu	4978	4978	1.028	19.52
Scotch	64	eur	6772	6772	1.031	77.41
kMetis	64	rgg20	42465	41066	1.030	1.58
kMetis	64	Delaunay20	28543	28318	1.030	1.21
kMetis	64	deu	5385	5147	1.029	5.31
kMetis	64	eur	12738	11313	1.070	30.30
parMetis	64	rgg20	43545	42863	1.050	0.55
parMetis	64	Delaunay20	30321	29535	1.047	0.65
parMetis	64	deu	7273	7083	1.027	0.91
parMetis	64	eur	16427	14976	1.025	5.65

Table 5: Performance for the largest graphs with coordinate information.

graph	avg. cut	best. cut.	avg. balance	avg. runtime
rgg20	15442	15039	1.029	7.20
Delaunay20	11533	11307	1.028	6.31
fetooth	19813	19559	1.029	0.65
598a	28596	27983	1.030	6.76
feocean	9553	9457	1.029	0.70
144	41977	40264	1.030	6.63
wave	47270	46293	1.029	1.47
m14b	49397	48769	1.030	6.41
auto	86001	84236	1.030	12.25
deu	1656	1593	1.029	21.79
eur	2048	1931	1.026	94.56
afshell10	175918	174677	1.029	17.00
coAuthorsDBLP	163463	161842	1.030	9.99
citationCiteseer	254914	253359	1.030	20.07

Table 6: KaPPa-Minimal $k = 16$

graph	avg. cut	best. cut.	avg. balance	avg. runtime
rgg20	24164	23842	1.029	3.94
Delaunay20	18179	17993	1.029	3.33
fetooth	28391	28070	1.030	0.53
598a	43741	43111	1.030	7.74
feocean	15657	15465	1.030	0.47
144	62171	61774	1.030	8.79
wave	68620	68085	1.030	1.02
m14b	73598	72484	1.030	8.12
auto	133723	131545	1.030	20.23
deu	2711	2626	1.029	11.50
eur	3386	3202	1.029	55.63
afshell10	275149	270249	1.030	9.25
coAuthorsDBLP	172830	171784	1.030	8.57
citationCiteseer	285710	278587	1.030	19.83

Table 7: KaPPa-Minimal $k = 32$

graph	avg. cut	best. cut.	avg. balance	avg. runtime
rgg20	35629	35252	1.030	2.09
Delaunay20	27001	26314	1.029	1.79
fetooth	39095	38423	1.029	0.62
598a	61924	61396	1.029	6.21
feocean	24275	24147	1.030	0.51
144	86950	86067	1.030	8.16
wave	93424	92366	1.030	1.03
m14b	107173	106361	1.030	10.24
auto	187424	185836	1.030	25.39
deu	4317	4193	1.029	5.97
eur	5770	5569	1.029	29.64
afshell10	404085	400378	1.030	4.82
coAuthorsDBLP	180724	180059	1.030	15.82
citationCiteseer	315062	313465	1.030	22.86

Table 8: KaPPa-Minimal $k = 64$

graph	avg. cut	best. cut.	avg. balance	avg. runtime
rgg20	15339	15013	1.029	24.61
Delaunay20	11061	10882	1.029	48.05
fetooth	18524	18198	1.030	3.55
598a	26887	26670	1.030	12.51
feocean	8469	8294	1.030	3.04
144	39492	39266	1.030	17.53
wave	45202	44936	1.030	10.73
m14b	46108	45931	1.030	19.27
auto	80683	79711	1.030	58.20
deu	1618	1556	1.027	78.82
eur	1935	1907	1.028	295.81
afshell10	166480	165625	1.030	69.97
coAuthorsDBLP	150272	149302	1.030	66.47
citationCiteseer	203302	198450	1.030	85.02

Table 9: KaPPa-Fast $k = 16$

graph	avg. cut	best. cut.	avg. balance	avg. runtime
rgg20	24222	23383	1.030	16.93
Delaunay20	17150	16814	1.030	24.44
fetooth	26677	26404	1.030	2.92
598a	41186	40928	1.030	11.91
feocean	14042	13618	1.030	2.15
144	58652	58175	1.030	16.03
wave	64532	64004	1.030	8.19
m14b	69223	68715	1.030	17.99
auto	125876	124920	1.030	46.44
deu	2641	2535	1.029	41.93
eur	3314	3231	1.030	306.52
afshell10	255746	252487	1.030	52.00
coAuthorsDBLP	163767	162577	1.030	58.93
citationCiteseer	233459	229629	1.030	83.14

Table 10: KaPPa-Fast $k = 32$

graph	avg. cut	best. cut.	avg. balance	avg. runtime
rgg20	35539	35086	1.030	9.95
Delaunay20	25129	24946	1.030	12.83
fetooth	36992	36795	1.029	2.57
598a	59233	59026	1.029	9.64
feocean	21973	21809	1.030	2.02
144	82493	82029	1.030	12.05
wave	89297	88924	1.030	6.09
m14b	101861	101410	1.030	17.46
auto	178119	177461	1.030	44.14
deu	4146	4078	1.029	31.63
eur	5538	5448	1.030	183.98
afshell10	384140	380225	1.030	29.43
coAuthorsDBLP	174411	173629	1.030	65.76
citationCiteseer	269854	268188	1.030	86.06

Table 11: KaPPa-Fast $k = 64$

graph	avg. cut	best. cut.	avg. balance	avg. runtime
rgg20	15199	14953	1.029	35.86
Delaunay20	11008	10816	1.027	67.92
fetooth	18570	18302	1.030	7.18
598a	26825	26467	1.030	17.74
feocean	8350	8188	1.030	5.62
144	39319	39010	1.030	26.04
wave	45048	44831	1.030	20.54
m14b	45762	45352	1.030	28.11
auto	79769	78713	1.030	87.41
deu	1616	1550	1.027	105.96
eur	1900	1760	1.027	497.93
afshell10	166427	165025	1.030	106.63
coAuthorsDBLP	145975	145031	1.030	105.61
citationCiteseer	176690	171233	1.030	142.01

Table 12: KaPPa-Strong $k = 16$

graph	avg. cut	best. cut.	avg. balance	avg. runtime
rgg20	23917	23430	1.029	26.04
Delaunay20	17086	16813	1.030	42.67
fetooth	26617	26397	1.030	5.28
598a	41190	40946	1.030	18.16
feocean	13815	13593	1.030	4.34
144	58631	58331	1.030	24.60
wave	64390	63981	1.030	14.94
m14b	69075	68107	1.030	29.94
auto	125500	124606	1.030	71.77
deu	2615	2548	1.029	73.17
eur	3291	3186	1.029	417.52
afshell10	255535	253525	1.030	80.85
coAuthorsDBLP	161073	160225	1.030	106.63
citationCiteseer	207559	203989	1.030	140.53

Table 13: KaPPa-Strong $k = 32$

graph	avg. cut	best. cut.	avg. balance	avg. runtime
rgg20	35354	34778	1.030	11.62
Delaunay20	25179	24799	1.030	22.04
fetooth	37002	36862	1.029	4.71
598a	59387	59148	1.029	14.15
feocean	21859	21636	1.030	3.68
144	82452	82286	1.030	19.11
wave	88964	88376	1.030	12.51
m14b	101455	101053	1.030	25.26
auto	177595	177038	1.030	62.64
deu	4093	4021	1.029	49.55
eur	5393	5290	1.030	308.17
afshell10	382923	379125	1.030	43.01
coAuthorsDBLP	172132	171194	1.030	111.90
citationCiteseer	249544	246150	1.030	146.65

Table 14: KaPPa-Strong $k = 64$

graph	avg. cut	best. cut.	avg. balance	avg. runtime
rgg20	18125	17498	1.021	1.53
Delaunay20	12440	11854	1.016	1.14
fetooth	20386	20035	1.029	0.09
598a	28854	27857	1.030	0.17
feocean	10377	10115	1.029	0.13
144	43041	42861	1.030	0.24
wave	49000	48404	1.030	0.22
m14b	49269	48314	1.029	0.36
auto	89139	85562	1.030	0.91
deu	2161	2041	1.007	5.19
eur	9395	3519	1.030	30.58
afshell10	188765	184350	1.014	3.06
coAuthorsDBLP	139658	138334	1.031	0.98
citationCiteseer	157011	153588	1.031	1.05

Table 15: KMetis $k = 16$

graph	avg. cut	best. cut.	avg. balance	avg. runtime
rgg20	18760	18193	1.048	0.39
Delaunay20	13126	12806	1.043	0.35
fetooth	20686	20255	1.046	0.06
598a	29858	29308	1.047	0.17
feocean	10212	9951	1.043	0.06
144	43019	41841	1.050	0.19
wave	49981	49537	1.048	0.09
m14b	49621	47697	1.048	0.28
auto	87057	84900	1.047	0.54
deu	3166	3063	1.009	1.62
eur	6861	5576	1.073	12.85
afshell10	191995	189925	1.048	0.74
coAuthorsDBLP	193580	190892	1.044	1.44
citationCiteseer	197095	197095	1.047	1.41

Table 16: parMetis $k = 16$

graph	avg. cut	best. cut.	avg. balance	avg. runtime
rgg20	28495	27765	1.029	1.58
Delaunay20	19304	18816	1.029	1.18
fetooth	29052	28547	1.030	0.10
598a	44213	43256	1.030	0.19
feocean	16877	16565	1.030	0.15
144	62481	61716	1.030	0.26
wave	68604	68062	1.030	0.25
m14b	74135	72746	1.030	0.40
auto	134086	133026	1.030	0.99
deu	3445	3319	1.019	5.28
eur	9442	7424	1.078	30.81
afshell10	291590	289400	1.027	3.13
coAuthorsDBLP	160373	159032	1.030	1.19
citationCiteseer	201073	197839	1.031	1.19

Table 17: KMetis $k = 32$

graph	avg. cut	best. cut.	avg. balance	avg. runtime
rgg20	29227	28650	1.049	0.22
Delaunay20	20141	19803	1.045	0.21
fetooth	28790	28513	1.043	0.07
598a	44422	43968	1.046	0.49
feocean	16259	16010	1.040	0.05
144	62673	62244	1.049	0.51
wave	70365	70072	1.048	0.15
m14b	76447	75356	1.049	0.52
auto	137913	137047	1.047	0.70
deu	4858	4703	1.034	0.87
eur	9616	8366	1.072	7.22
afshell10	293110	289275	1.048	0.35
coAuthorsDBLP	211756	209846	1.046	1.59
citationCiteseer	212524	212524	1.050	1.56

Table 18: parMetis $k = 32$

graph	avg. cut	best. cut.	avg. balance	avg. runtime
rgg20	42465	41066	1.030	1.58
Delaunay20	28543	28318	1.030	1.21
fetooth	39381	39233	1.030	0.12
598a	62703	61888	1.030	0.22
feocean	24531	24198	1.030	0.17
144	87208	86534	1.030	0.30
wave	94083	92148	1.030	0.29
m14b	108141	107384	1.031	0.44
auto	189699	188555	1.030	1.08
deu	5385	5147	1.029	5.31
eur	12738	11313	1.070	30.30
afshell10	427047	421285	1.030	3.18
coAuthorsDBLP	176485	174402	1.033	1.42
citationCiteseer	244330	242677	1.033	1.41

Table 19: KMetis $k = 64$

graph	avg. cut	best. cut.	avg. balance	avg. runtime
rgg20	43545	42863	1.050	0.55
Delaunay20	30321	29535	1.047	0.65
fetooth	39477	38790	1.047	0.56
598a	63688	62936	1.047	1.82
feocean	26249	25912	1.039	0.12
144	87967	87163	1.047	1.58
wave	95758	94605	1.049	0.44
m14b	108546	107125	1.049	1.98
auto	194958	192198	1.047	1.69
deu	7273	7083	1.027	0.91
eur	16427	14976	1.025	5.65
afshell10	435995	433525	1.049	0.20
coAuthorsDBLP	218798	217403	1.050	2.32
citationCiteseer	219850	219850	1.046	2.32

Table 20: parMetis $k = 64$

Graph	2		4		8		16		32		64	
3elt	** 90	89	+ 201	199	* 354	342	* 597	569	* 1008	969	* 1629	1564
add20	* 618	594	* 1190	1177	* 1752	1704	+ 2141	2121	* 2594	2687	* 3082	3236
data	+ 191	188	* 383	383	* 664	660	** 1169	1162	* 1912	1865	* 2949	2885
uk	* 20	19	+ 44	42	* 88	84	+ 159	152	* 273	258	** 445	438
add32	** 10	10	** 33	33	** 66	66	+ 124	117	+ 223	212	* 495	720
bcsstk33	** 10169	10097	* 21800	21508	** 34560	34178	* 56639	54860	* 80237	78132	+ 111075	108505
whitaker3	* 127	126	* 383	380	+ 668	656	** 1150	1093	** 1754	1717	** 2676	2567
crack	** 184	183	* 370	362	* 694	678	** 1160	1092	+ 1815	1707	+ 2717	2566
wingnodal	* 1710	1696	** 3626	3572	** 5588	5443	** 8566	8422	* 12384	11980	+ 16716	16134
fe4elt2	** 130	130	* 349	349	+ 616	605	** 1032	1014	+ 1694	1657	* 2640	2537
vibrobox	* 11308	10310	+ 19249	19199	+ 24923	24553	+ 34505	32167	** 42432	41399	** 51229	49521
bcsstk29	** 2853	2818	** 8156	8379	* 14813	13965	* 23914	21768	* 37309	34886	+ 58987	57054
4elt	** 139	138	** 329	321	** 555	534	** 989	939	* 1639	1559	** 2718	2596
fesphere	** 386	386	* 794	768	** 1215	1152	* 1881	1730	* 2745	2565	+ 3968	3663
cti	** 334	318	* 973	944	* 1836	1802	** 2990	2906	* 4375	4223	* 6346	5875
memplus	** 5712	5489	* 9562	9584	** 12190	11785	* 13908	13241	* 15587	14489	+ 17381	17063
cs4	+ 389	367	* 1003	940	+ 1568	1470	** 2302	2206	* 3228	3090	* 4458	4169
bcsstk30	** 6391	6335	+ 16651	16622	* 35037	34604	* 73118	71234	* 119316	115770	* 180243	173945
bcsstk31	+ 2769	2701	* 7512	7444	** 13608	13417	* 24821	24277	** 39455	38086	** 61327	60528
fepwt	** 342	340	** 712	705	** 1454	1442	* 2844	2806	* 5637	5758	** 8648	8454
bcsstk32	* 4667	4667	* 9440	9538	* 21800	21490	** 37701	37673	* 63382	61144	* 98842	95199
febbody	+ 266	262	* 649	671	* 1100	1156	* 1910	1931	* 3106	3202	* 5212	5282
t60k	* 84	75	* 220	211	* 483	465	+ 891	849	* 1466	1391	** 2297	2211
wing	* 851	787	* 1793	1666	* 2720	2589	* 4203	4131	* 6217	5902	+ 8534	8132
brack2	** 731	708	* 3121	3038	* 7363	7269	** 12177	11983	** 18236	17798	* 27442	26557
finan512	** 162	162	* 324	324	* 648	648	** 1296	1296	* 2592	2592	** 10862	10560
fetooth	* 3893	3823	* 7096	7103	* 11953	12060	+ 18227	18283	* 26517	25977	+ 37079	35980
ferotor	+ 2103	2045	** 7461	7694	** 13283	13165	+ 21249	20773	** 33266	32783	** 49079	47461
598a	* 2426	2388	* 8131	8197	* 16491	16594	* 26838	27009	** 40471	40962	* 59445	59098
feocean	** 468	387	* 1914	1878	+ 4270	4538	* 8447	8507	+ 13673	13767	** 21774	21854
144	* 6604	6479	* 16162	15345	** 26266	25818	* 39195	39352	* 58702	58126	* 82904	81145
wave	* 8812	8682	** 17616	17950	+ 30375	31697	+ 44783	44711	+ 64646	64860	* 89332	88863
m14b	* 3871	3826	** 13296	13403	* 26657	27066	* 44013	44330	* 69072	67770	+ 102393	101551
auto	+ 10329	10042	* 28051	27790	* 47321	48442	* 79741	81339	* 126146	124991	** 179095	175975

Table 21: Walshaw Benchmark with $\epsilon = 1\%$. * Expansion*, ** Expansion*², + InnerOuter.

Graph	2		4		8		16		32		64	
3elt	** 87	87	+ 200	198	* 343	336	+ 584	565	* 1010	958	+ 1622	1542
add20	+ 619	576	* 1179	1158	** 1790	1690	+ 2161	2095	** 2559	2493	+ 3058	3152
data	+ 193	185	** 380	378	+ 665	650	** 1157	1133	** 1912	1802	* 2936	2809
uk	** 18	18	+ 42	40	** 82	81	* 151	148	* 265	251	** 440	414
add32	+ 10	10	** 33	33	** 66	66	+ 124	117	+ 222	212	* 494	624
bcsstk33	+ 10064	10064	** 21195	21035	+ 34386	34078	** 56262	54510	** 80001	77672	+ 110822	107012
whitaker3	+ 126	126	** 384	378	** 665	655	+ 1138	1092	* 1753	1686	+ 2655	2535
crack	+ 182	182	* 360	360	* 678	676	+ 1126	1082	* 1782	1679	+ 2670	2553
wingnodal	** 1682	1680	* 3565	3566	+ 5430	5401	+ 8451	8316	* 12277	11938	+ 16702	15971
fe4elt2	+ 130	130	+ 349	343	** 608	598	** 1015	1007	** 1681	1633	** 2617	2527
vibrobox	** 11188	10310	** 19107	18778	** 24531	24171	+ 34189	31516	* 42650	39592	+ 50183	49123
bcsstk29	+ 2818	2818	* 8153	8045	+ 14437	13817	+ 23532	21410	** 37015	34407	+ 58738	55366
4elt	+ 138	137	** 320	319	+ 536	523	+ 953	914	* 1624	1537	+ 2715	2581
fesphere	+ 384	384	* 796	764	+ 1217	1152	* 1851	1706	* 2719	2477	* 3767	3547
cti	+ 318	318	** 927	917	* 1773	1716	* 2895	2778	* 4263	4132	* 6207	5763
memplus	+ 5532	5355	* 9953	9418	+ 12239	11628	+ 13755	13237	* 15432	14350	+ 17870	17002
cs4	+ 383	362	* 1001	936	** 1542	1470	+ 2237	2126	+ 3164	3048	* 4397	4169
bcsstk30	+ 6251	6251	* 16528	16577	** 34505	34559	* 72618	70278	* 118106	114005	+ 179278	171727
bcsstk31	+ 2676	2676	** 7209	7258	* 13253	13246	* 24365	23504	* 38817	37459	** 60577	58667
fepwt	+ 340	340	+ 705	704	+ 1418	1421	* 2789	2784	+ 5603	5606	** 8630	8346
bcsstk32	+ 4667	4667	+ 8805	9533	+ 20992	21307	+ 36628	37204	* 62639	59824	** 97535	92690
febbody	+ 265	262	* 613	668	* 1055	1094	* 1798	1903	+ 2928	3086	* 4997	5212
t60k	+ 74	71	* 211	207	* 470	454	* 875	822	+ 1443	1391	** 2272	2198
wing	** 840	774	* 1761	1636	* 2661	2551	** 4144	4015	* 6107	5832	** 8340	8043
brack2	685	684	* 2840	2864	* 7105	6994	* 11687	11741	* 17815	17649	+ 26755	26366
finan512	+ 162	162	+ 324	324	+ 648	648	* 1296	1296	* 2592	2592	** 10944	10560
fetooth	* 3807	3792	+ 6947	7081	* 11562	11957	* 17678	18093	* 25884	25624	* 36178	35830
ferotor	** 1964	1965	+ 7263	7636	** 12798	12862	+ 20404	20521	+ 32155	31763	* 47808	47049
598a	* 2373	2367	* 7963	7978	* 16079	16031	* 25960	26257	** 39792	40718	58430	58454
feocean	+ 311	311	+ 1706	1704	* 3976	4019	** 8004	7838	** 13196	12746	* 21060	21854
144	+ 6512	6438	+ 15555	15250	** 25529	25611	** 38701	38478	** 57561	57354	* 80981	80767
wave	* 8699	8616	* 16947	17407	** 29022	29776	* 43168	43791	+ 62766	63675	+ 87272	87957
m14b	* 3833	3823	* 13131	13285	* 26044	26153	* 42942	43962	* 67272	67551	+ 100112	101019
auto	** 9806	9782	+ 26343	26509	** 45703	48263	** 77461	80495	* 123442	124251	** 175520	174904

Table 22: Walsh Benchmark with $\epsilon = 3\%$. * Expansion*, ** Expansion *2 , + InnerOuter.

Graph	2		4		8		16		32		64	
3elt	** 87	87	** 199	197	+ 339	330	+ 581	560	** 1001	950	** 1615	1539
add20	** 579	550	** 1179	1157	+ 1744	1675	+ 2150	2081	* 2560	2463	+ 3054	3152
data	* 188	181	** 374	368	** 650	628	+ 1147	1086	* 1888	1777	+ 2910	2798
uk	** 18	18	+ 41	40	** 81	78	+ 152	139	** 262	246	+ 437	410
add32	** 10	10	** 33	33	** 66	65	+ 124	117	+ 222	212	** 494	624
bcsstk33	** 9914	9914	* 20614	20584	+ 34190	33938	** 55868	54323	+ 79530	77163	* 110822	106886
whitaker3	** 126	126	** 382	378	** 665	650	* 1130	1084	* 1737	1686	+ 2655	2535
crack	** 182	182	** 360	360	** 679	667	+ 1122	1080	* 1755	1679	+ 2651	2548
wingnodal	* 1676	1668	* 3545	3536	+ 5376	5350	+ 8388	8316	** 12252	11879	** 16595	15873
fe4elt2	** 130	130	** 349	335	* 599	583	** 1015	991	** 1660	1633	+ 2609	2516
vibrobox	** 11188	10310	** 18958	18778	* 24121	23930	* 33760	31235	** 42269	39592	+ 49552	48200
bcsstk29	** 2818	2818	+ 8055	7942	* 14009	13614	+ 23131	20924	** 36633	33818	* 58183	54935
4elt	** 137	137	* 319	315	* 526	516	** 946	902	* 1590	1532	* 2675	2565
fesphere	** 384	384	** 784	764	+ 1217	1152	* 1840	1692	** 2709	2477	+ 3945	3547
cti	** 318	318	+ 891	897	* 1737	1716	** 2885	2725	* 4242	4037	+ 6010	5684
memplus	** 5528	5267	+ 9489	9299	** 12091	11555	+ 13701	13078	** 15362	14249	+ 17632	16662
cs4	* 373	356	* 990	936	** 1542	1470	+ 2237	2126	* 3141	2995	** 4364	4116
bcsstk30	** 6251	6251	** 16316	16417	** 34391	34559	+ 72087	70043	** 117512	113321	** 177303	170591
bcsstk31	** 2676	2676	** 7118	7223	* 13104	13058	+ 24062	23254	** 38279	37459	** 60257	57534
fepwt	** 340	340	** 700	704	** 1406	1411	+ 2773	2778	+ 5525	5606	+ 8582	8310
bcsstk32	** 4667	4667	+ 8539	9052	** 20568	20099	+ 35962	35990	** 61021	59824	** 96032	91006
febbody	** 263	262	* 599	629	* 1055	1072	* 1786	1815	+ 2863	3086	+ 4897	5043
t60k	** 69	65	+ 206	196	* 469	454	* 865	818	** 1436	1376	** 2263	2168
wing	+ 826	770	** 1734	1636	* 2632	2551	* 4106	4015	* 6063	5806	+ 8300	7991
brack2	** 660	660	** 2739	2755	* 6776	6883	+ 11557	11558	* 17617	17529	+ 26555	26281
finan512	** 162	162	** 324	324	** 648	648	* 1296	1296	** 2592	2592	** 10909	10560
fetooth	** 3785	3773	* 6863	7027	+ 11498	11957	* 17509	18090	* 25641	25624	+ 35795	35476
ferotor	** 1955	1957	+ 7031	7520	* 12643	12678	** 20098	20263	+ 31611	31576	+ 47186	46608
598a	** 2344	2336	+ 7837	7978	** 15794	16031	* 25782	26257	** 39478	40179	** 58180	58307
feocean	** 311	311	** 1688	1704	* 3952	4019	+ 7671	7838	+ 12953	12746	+ 20660	21784
144	* 6502	6362	+ 15313	15250	** 25529	25611	+ 38182	38478	+ 57202	57354	+ 80653	80257
wave	** 8613	8563	+ 16780	17306	+ 28753	29776	+ 42810	43791	** 62382	63675	** 86867	87654
m14b	** 3844	3802	* 13124	13285	** 25701	26153	+ 42644	43747	* 66845	67551	+ 99460	100183
auto	* 9587	9450	+ 25805	26097	** 44915	48174	+ 76500	80495	** 121988	124251	** 174173	174904

Table 23: Walsh Benchmark with $\epsilon = 5\%$. * Expansion*, ** Expansion *2 , + InnerOuter.